


Search trajectory: proposal quality before and after a logged failure mode was corrected



Both later runs start from the same converged graph (0.602689). The per-iteration logs made the failure mode explicit: every final_05 proposal replaced the incumbent model outright rather than building on it. Rewriting the proposer's search-strategy guidance between runs made all four subsequent proposals additive and produced a new best of 0.603138 — +0.001669 over the official baseline, reached in four iterations with zero manual interventions inside any scored run.